

Buffer Overflow Attack



Lecture 2

Instructor: C. Pu (Ph.D., Assistant Professor)

puc@marshall.edu



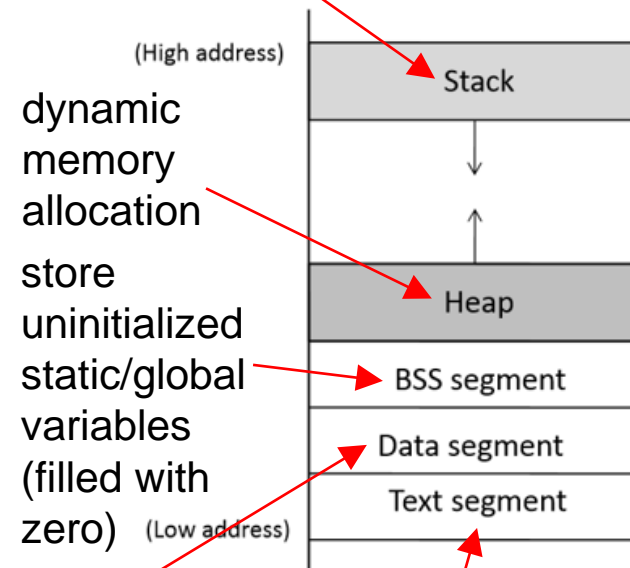
Introduction

- famous buffer overflow attacks
 - Morris worm (1988)
 - buffer overflow in the fingerd network service
 - Code Red worm (2001)
 - execute arbitrary code and infect the machine with the worm
 - SQL Slammer (2003)
 - generate random IP addresses and send itself out to those addresses
 - Stagefright attack against Android (2015)
 - allows adversary to perform arbitrary operations on the victim's device
 - more...

Program Memory Layout

- prerequisite of understanding buffer overflow attack:
 - understanding how the data memory is arranged inside a process
- when program running, needs memory space to store data
 - for C program, its memory is divided into five segments
 - text segment
 - data segment
 - BSS segment
 - heap
 - stack

store local variables defined inside functions, and function-related data (return address)



store static/global variables

store executable code of program (read-only)

Program Memory Layout (cont.)

allocates size bytes of uninitialized storage

! arg: number of bytes to allocate

ref: <https://en.cppreference.com/w/c/memory/malloc>

```
int x = 100;
int main()
{
    // data stored on stack
    int a=2;
    float b=2.5;
    static int y;

    // allocate memory on heap
    int *ptr = (int *) malloc(2*sizeof(int));

    // values 5 and 6 stored on heap
    ptr[0]=5;
    ptr[1]=6;

    // deallocate memory on heap
    free(ptr);

    return 1;
}
```

pointer
variable

int pointer

return the size
of data type

a, b, ptr


ptr points to the
memory here

y

x


(High address)

(Low address)



store executable
code of program
(read-only)

Stack Memory Layout



- stack: store data used in function invocations
- a program executes as a series of function calls (execution)
 - when function is called, space is allocated for it on the stack
 - e.g.,

```
void func(int a, int b)
{
    int x, y;

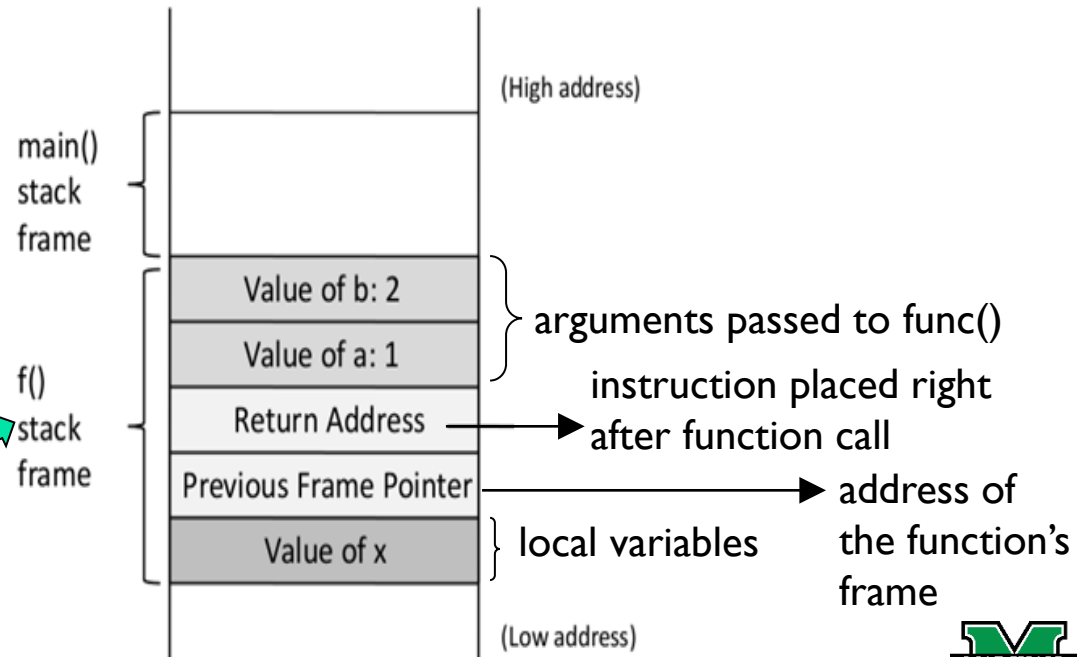
    x = a + b;
    y = a - b;
}
```

two integer arguments: a and b

two integer local variables: x and y

when func() is called, stack frame is allocated

Stack grows



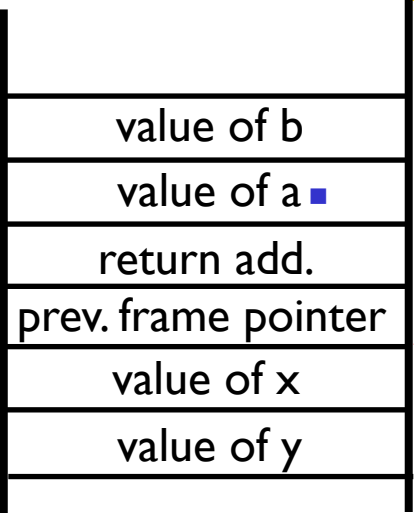
Frame Pointer

```
void func(int a, int b)
{
    int x, y;
    store result get value
    x = a + b;
    y = a - b;
}
```

- inside func(), how to access arguments and local variables?
 - only way: knowing their memory add.
 - issue: add. cannot be determined during compilation (compilers cannot predict run-time status of stack)
 - solution: frame pointer, a special register in CPU
 - points to a fixed location in stack frame
 - the add. of each argument and local variable can be calculated using frame pointer and add. offset

- the value of offset can be decided during compilation

the value of frame pointer can change during run time



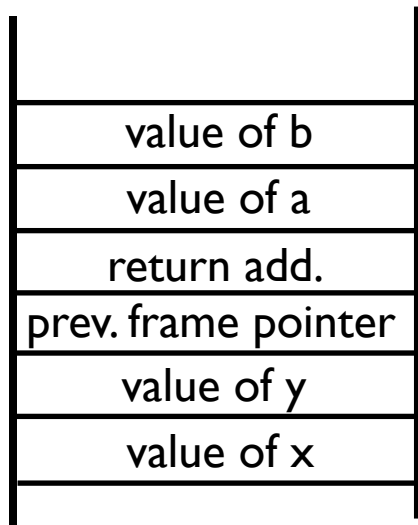
current
frame
pointer

```
movl    12(%ebp), %eax    ; b is stored in %ebp + 12
movl    8(%ebp), %edx     ; a is stored in %ebp + 8
addl    %edx, %eax
movl    %eax, -8(%ebp)    ; x is stored in %ebp - 8
```

Frame Pointer

```
void func(int a, int b)
{
    int x, y;

    x = a + b;
    y = a - b;
}
```



current
frame
pointer

on 32-bit architecture,
return address and frame
pointer both occupy 4 bytes.
So,
a is at $\text{ebp} + 8$
b is at $\text{ebp} + 12$

frame pointer register (x86 architecture)

```
movl    12(%ebp), %eax    ; b is stored in %ebp + 12
movl    8(%ebp), %edx     ; a is stored in %ebp + 8
addl    %edx, %eax
movl    %eax, -8(%ebp)    ; x is stored in %ebp - 8
```

eax and edx: general-purpose registers storing temporary values

$12(\text{ebp}): \text{ebp} + 12$

`movl array_base(%esi), %eax`

add the address of memory location `array_base` to the contents of number register `%esi` to determine an address in memory. Then move the contents of this address into number register `%eax`.

`addl %edx, %eax`


adds together its two operands (`%edx` and `%eax`), storing the result in its second operand (`%eax`)

$-8(\text{ebp}): \text{ebp} - 8$

Function Call Chain


- call another function from inside a function
 - every time function is called, a stack frame is allocated on the top of stack
 - when function is returned (completed), the stack frame allocated for it is released
 - e.g.,

```
void f(int a, int b)
{
    int x;
}
void main()
{
    f(1,2);
    printf("hello world");
}
```



Function Call Chain (cont.)

- only one frame pointer register: always pointing to the frame of current function
- question: how the functions were called?





Stack Buffer-Overflow Attack

- memory copying: copying data from one place to another place
 - before copying, a program allocates memory space
 - issue: programmer fails to allocate *sufficient* amount of memory
 - consequence: more data is copied to the des. buffer than the amount of allocated space ➡ buffer overflow
 - program crash (corruption of data beyond buffer)
 - gain control of program (attacker)
- some languages (e.g., Java) automatically detect the problem (buffer over-run), while many others (e.g., C) do not



Copying Data Causes Buffer Overflow

■ strcpy()

```
#include <string.h>
#include <stdio.h>
```

```
void main()
```

```
{
```

```
    char src[40] = "hello world \0 extra string";
```

```
    char dest[40];
```

```
    // copy to dest (destination) from src (source)
```

```
    strcpy(dest, src); → only copy "hello world" to dest. why???
```

```
}
```

char* strcpy(char* destination, const char* source):

- copies the string pointed by the source (including the null character) to the destination.
- when making copy, it stops when meets \0 (the end of string)

Copying Data Causes Buffer Overflow

- when copying data, what will happen if the string is longer than the size of buffer?

```
#include <string.h>
```

```
void foo(char *str)
```

```
{
```

```
    char buffer[12];
```

```
    strcpy(buffer, str);
```

```
}
```


```
void main()
```

```
{
```

```
    char *str = "This is definitely longer than 12";
```

```
    foo(str);
```

```
}
```





Exploiting Buffer Overflow Vulnerability

- overflowing buffer:
 - cause program crash
 - run some other code (more interesting to attacker)
 - if attackers control what code to run, they can hijack the execution of programs
 - privilege escalation for attackers

Exploiting Buffer Overflow Vulnerability (cont.)


```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
int foo(char *str){
    char buffer[100];
    /* The following statement has a buffer overflow problem */
    strcpy(buffer, str);
    return 1;
}
```

```
void main(int argc, char **argv){
    char str[400];
    FILE *badfile;
    badfile = fopen("badfile", "r");
    fread(str, sizeof(char), 300, badfile);
    foo(str);
    printf("Returned Properly\n");
}
```


open file "badfile" to read.

read 300 bytes and copy
data to 100 bytes buffer
the content is copied to
buffer from "badfile"

do you know
what inside?



Exploiting Buffer Overflow Vulnerability (cont.)

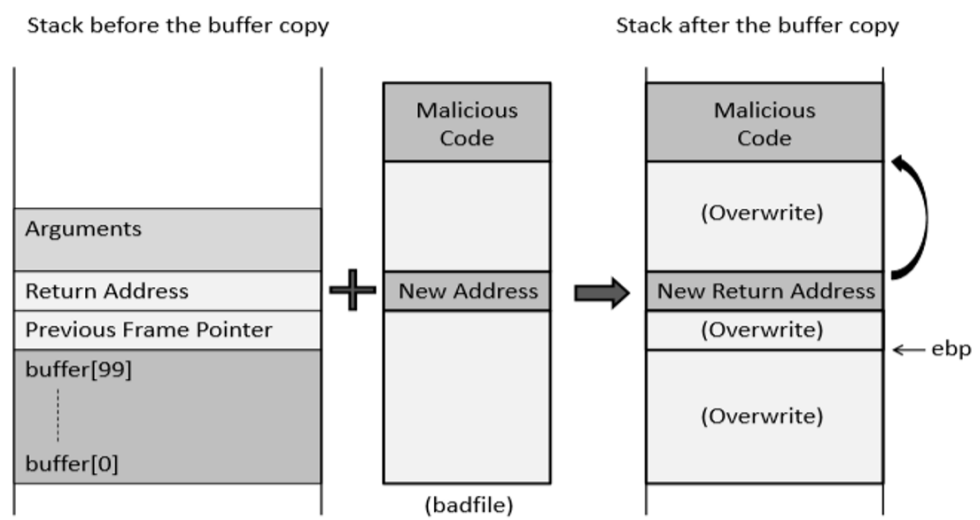


get code into memory of running program:

- not difficult:
 - place code in “badfile”
 - let program read “badfile”
 - program copies code to buffer

force program to jump to our code (already in memory)

- using buffer overflow
 - overwrite return add. field
 - use add. of malicious code to overwrite
 - when foo() returns, it jumps to new add. (add. of malicious code)





Setup for Environment

- attack environment: Ubuntu
- buffer overflow has a long history, so many OS have countermeasures against it
- to simplify environment
 - turn off countermeasures
 - later on, turn them back on to show
 - countermeasures only make buffer overflow more difficult, not impossible



Disable Address Randomization

- address space layout randomization (ASLR):
countermeasure to buffer overflow
 - randomizing the memory space of key data areas in process
 - the base of executable
 - the positions of stack, heap, and libraries
 - making it difficult for attackers to guess the add. of injected malicious code



Disable Address Randomization

- turning countermeasure off

```
% sudo sysctl -w kernel.randomize_va_space=0
```

- goal: exploit buffer overflow vulnerability in Set-UID root program
 - a Set-UID root program runs with root privilege when executed by normal user
 - assigning normal user extra privileges
 - if buffer overflow vulnerability is exploited in privileged Set-UID root program
 - consequence: the injected malicious code can run with root's privilege

Vulnerable Program: stack.c

- compile set-uid root version of program

```
% gcc -o stack -z execstack -fno-stack-protector stack.c
% sudo chown root stack
% sudo chmod 4755 stack
```

make stack
executable

turn off Stack-Guard
(countermeasure)

- 1st command: compiles stack.c program
- 2nd and 3rd commands: turn executable stack into root-owned set-uid program
 - the order of 2nd and 3rd commands cannot be reversed

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>

int foo(char *str){
    char buffer[100];
    /* The following statement has a buffer
    overflow problem */
    strcpy(buffer, str);
    return 1;
}

void main(int argc, char **argv){
    char str[400];
    FILE *badfile;
    badfile = fopen("badfile", "r");
    fread(str, sizeof(char), 300, badfile);
    foo(str);
    printf("Returned Properly\n");
}
```



Vulnerable Program: stack.c

- badfile: contains random contents
 - when the size of file is less than 100 bytes, the program runs properly
 - when the size of file is larger than 100 bytes, the program crashes
 - buffer overflow happens

```
$ echo "aaaa" > badfile
```

```
$ ./stack
```

```
Returned Properly
```

```
$
```

```
$ echo "aaa ... (100 characters omitted) ... aaa" > badfile
```

```
$ ./stack
```

```
Segmentation fault (core dumped)
```