

Buffer Overflow Countermeasures



Lecture 3

Instructor: C. Pu (Ph.D., Assistant Professor)

puc@marshall.edu



Introduction



- countermeasures proposed and deployed in real-world systems and software
 - from hardware architecture, OS, compiler, library, to applications



Safer Functions

- to some memory copy functions, certain special characters decide whether the copy should end or not
 - *dangerous*: the length of data that can be copied is decided by data, which may be controlled by users
 - e.g., *strcpy*, *sprintf*, *strcat*, and *gets*
 - *strcpy*: <https://www.cplusplus.com/reference/cstring/strcpy/>
 - *sprintf*: <https://www.cplusplus.com/reference/cstdio/sprintf/?kw=sprintf>
 - *strcat*: <https://www.cplusplus.com/reference/cstring/strcat/?kw=strcat>
 - *gets*: <https://www.cplusplus.com/reference/cstdio/gets/?kw=gets>



Safer Functions (cont.)

- *safer* approach:
 - let developer have control: specifying the length in code
 - the size of target buffer decides the length, not the data
 - e.g., *strncpy*, *snprintf*, *strncat*, *fgets*
 - *strncpy*: <https://www.cplusplus.com/reference/cstring/strncpy/?kw=strncpy>
 - *snprintf*: <https://www.cplusplus.com/reference/cstdio/snprintf/?kw=snprintf>
 - *strncat*: <https://www.cplusplus.com/reference/cstring/strncat/?kw=strncat>
 - *fgets*: <https://www.cplusplus.com/reference/cstdio/fgets/?kw=fgets>
 - developers explicitly specify the max length of data to be copied into the target buffer
 - need to think about buffer size
 - relatively safer
 - what if developer intentionally specifies longer data?



Safer Dynamic Link Library

- drawback of safer function approach: require change in code
 - what if you only have binary?
 - difficult to change binary
- solution: dynamic link libraries (program uses library)
 - the library function code is not included in code's binary, instead, it is dynamically linked to code
 - safer library → safer code
- e.g., libsafe
 - perform boundary checking based on frame pointer
 - not allow copy beyond frame pointer
- e.g., libmib
 - support “limitless” strings, instead of fixed length string buffer
 - its own version functions like *strcpy*



Program Static Analyzer

- instead of eliminating buffer overflow, warns developers of buffer overflow vulnerabilities in code
 - implemented as command-line tool or in the editor
 - notify developers of unsafe code during developing phase
- e.g.,
 - ITS4 (C/C++)



Programing Language

- developer relies on programming language to develop program
- burden is removed if language does checking against buffer overflow
- e.g.,
 - Java and Python
 - automatic boundary checking
 - safer languages for avoiding buffer overflow



Other Approaches

- compiler: compile code, verify stack, and eliminate buffer overflow conditions
 - e.g., Stackshield and StackGuard:
 - check whether the return addr. has been modified before a function returns
 - Stackshield
 - idea: save a copy of return addr. at safer place
 - at beginning of function, compiler inserts instructions to copy return addr. to a location that cannot be overflowed
 - before returning from function, comparing return addr. on stack with the one in safer place
 - determine buffer overflow



Other Approaches

- compiler: compile code, verify stack, and eliminate buffer overflow conditions
 - e.g., Stackshield and StackGuard:
 - check whether the return addr. has been modified before a function returns
 - StackGuard
 - idea: put a guard between return addr. and buffer
 - if return addr. is modified, the guard will also be modified
 - at the start of function,
 - the compiler adds a random value below return addr.
 - save a copy of random value at safer location (off stack)
 - before function returns
 - the canary is checked against the saved value



Other Approaches (cont.)

- operating system
 - loader program:
 - before execution, program is loaded into system
 - running environment is set up
 - dictate how the memory of program is laid out
 - e.g., Address Space Layout Randomization (ASLR)
 - randomize the layout of program memory, making it hard for attacker to guess memory address
- hardware architecture
 - modern CPU supports NX bit (No-eXecute)
 - separate code from data
 - OS marks certain memory areas as non-executable
 - processor refuses to execute any code residing in these areas
 - if stack is marked as non-executable



Address Randomization

- to succeed in buffer overflow, attackers need to get vulnerable program to “return” to their injected code
 - guess where the injected code will be
 - (easy) predict where the stack is located in memory
 - most OS places stack in fixed location
- necessary to place stack in fixed location? No.
- when compiler generates binary code from source code
 - add. of data are not hard-coded in binary code
 - instead, their addr. are calculated based on frame and stack pointers
 - add. of data are represented as the offset of these two pointers



Address Randomization

- for attackers, they need to know the absolute addr., not the offset
 - important: knowing the stack location
- idea to defend against buffer overflow?
 - randomize the start location of stack
 - benefits:
 - make attacker's job more hard
 - no effect on program
- Address Layout Randomization (ASLR)



Address Randomization on Linux

- to run program, OS loads program into system
 - set up stack and heap memory for program
- memory randomization is normally implemented in loader
- for Linux, ELF is common binary format for program
 - randomization can be carried out by ELF loader



Address Randomization on Linux

- e.g., simple program with two buffers: stack and heap
 - print out their add. to see whether stack and heap are allocated in different places every time we run program

```
#include <stdio.h>
#include <stdlib.h>

void main() {
    char x[12];
    char *y = malloc (sizeof(char)*12);

    printf("address of buffer x (on stack): 0x%x\n", x);
    printf("address of buffer y (on heap): 0x%x\n", y);
}
```



Address Randomization on Linux

- e.g., simple program with two buffers: stack and heap
 - compile and run code under different randomization settings

```
// turn off randomization
$ sudo sysctl -w kernel.randomize_va_space=0
kernel.randomize_va_space =0
$ program.out
...
...
$ program.out
...
...
```



Address Randomization on Linux

- e.g., simple program with two buffers: stack and heap
 - compile and run code under different randomization settings

```
// randomize stack address
$ sudo sysctl -w kernel.randomize_va_space=1
kernel.randomize_va_space = 1
$ program.out
...
...
$ program.out
...
...
```




Address Randomization on Linux

- e.g., simple program with two buffers: stack and heap
 - compile and run code under different randomization settings

```
// randomize heap address
$ sudo sysctl -w kernel.randomize_va_space=2
kernel.randomize_va_space =2
$ program.out
...
...
$ program.out
...
...
```



StackGuard

- stack-based buffer overflow attack needs to modify return add.
- solution to stack-based buffer overflow attack:
 - detect whether the return add. is modified before returning from a function
- StackGuard: place a guard between return add. and buffer, and use the guard to detect whether the return add. is modified
 - incorporated into compilers like gcc

StackGuard

- place non-predictable value (called guard) between buffer and return add.

- before returning from function, check whether the value is modified
 - if modified, the return add. has been modified

detecting whether the return add. is overwritten

=

detecting whether the guard is modified

